

Learned vs random, per complexity (final)
SAC (target-entropy 0.4), 450 train / 207 held-out

